

# Used Cars Price Prediction

## Problem Definition

- There is a huge demand for used cars in the Indian Market today. As sales of new cars have slowed down in the recent past, the pre-owned car market has continued to grow over the past years and is larger than the new car market now. Cars4U is a budding tech start-up that aims to find footholes in this market.
- In 2018-19, while new car sales were recorded at 3.6 million units, around 4 million second-hand cars were bought and sold. There is a slowdown in new car sales and that could mean that the demand is shifting towards the pre-owned market. In fact, some car owners replace their old cars with pre-owned cars instead of buying new ones. The used car market is a very different beast with huge uncertainty in both pricing and supply. Keeping this in mind, the pricing scheme of these used cars becomes important in order to grow in the market.

## The Context:

- The used car market in India is larger than the new car market and is experiencing significant growth despite a slowdown in new car sales. This shift highlights the increasing consumer preference for pre-owned vehicles. Understanding and addressing the pricing scheme for these used cars is critical for companies like Cars4U to capitalize on the market potential and gain a competitive edge.

## The objective:

- The intended goal is to develop an accurate and reliable pricing model for used cars that reflects the current market dynamics. This model should help Cars4U offer competitive prices to both buyers and sellers, thereby increasing their market share and customer satisfaction.

## The key questions:

- Which features (e.g., make, model, age, mileage, location, etc.) significantly affect the price?
- What are the distributions of the variables?
- Are variables correlated?
- Which type of model is best suit for the task?

## The problem formulation:

- Develop a predictive pricing model: Create a machine learning model that predicts the price of a used car based on various attributes such as make, model, age, mileage, condition, location, and other relevant factors.

## Data Dictionary

**S.No.** : Serial Number

**Name** : Name of the car which includes Brand name and Model name

**Location** : The location in which the car is being sold or is available for purchase (Cities)

**Year** : Manufacturing year of the car

**Kilometers\_driven** : The total kilometers driven in the car by the previous owner(s) in KM

**Fuel\_Type** : The type of fuel used by the car (Petrol, Diesel, Electric, CNG, LPG)

**Transmission** : The type of transmission used by the car (Automatic / Manual)

**Owner** : Type of ownership

**Mileage** : The standard mileage offered by the car company in kmpl or km/kg

**Engine** : The displacement volume of the engine in CC

**Power** : The maximum power of the engine in bhp

**Seats** : The number of seats in the car

**New\_Price** : The price of a new car of the same model in INR 100,000

**Price** : The price of the used car in INR 100,000 (**Target Variable**)

## Loading libraries

```
In [79]: import numpy as np

import pandas as pd

import matplotlib.pyplot as plt

import seaborn as sns

from statsmodels.formula.api import ols

import statsmodels.api as sm

from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
from sklearn.metrics import r2_score, mean_absolute_percentage_error, mean_a

from sklearn.model_selection import train_test_split, GridSearchCV

from statsmodels.stats.diagnostic import het_white

from statsmodels.compat import lzip

import statsmodels.stats.api as sms

import pylab

import scipy.stats as stats

from sklearn.model_selection import KFold

from sklearn.linear_model import Ridge

from sklearn.tree import DecisionTreeRegressor

from sklearn.ensemble import RandomForestRegressor

import sklearn

import warnings
warnings.filterwarnings("ignore")
```

```
In [80]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

## Let us load the data

```
In [81]: data = pd.read_csv('/content/drive/MyDrive/MIT course/Capstone Project/used_
```

## Data Overview

- Observations
- Sanity checks

```
In [82]: data.head()
```

Out [82]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	C
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	

We can see all the information about a vehicle, there are missing values in 'New\_price' variable.

In [83]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   S.No.                 7253 non-null  int64
1   Name                  7253 non-null  object
2   Location              7253 non-null  object
3   Year                  7253 non-null  int64
4   Kilometers_Driven     7253 non-null  int64
5   Fuel_Type             7253 non-null  object
6   Transmission          7253 non-null  object
7   Owner_Type            7253 non-null  object
8   Mileage               7251 non-null  float64
9   Engine                7207 non-null  float64
10  Power                 7078 non-null  float64
11  Seats                 7200 non-null  float64
12  New_price             1006 non-null  float64
13  Price                 6019 non-null  float64
dtypes: float64(6), int64(3), object(5)
memory usage: 793.4+ KB
```

There are 14 variables total, 9 numerical, 5 categorical.

In [84]: `data.isnull().sum()`

```
Out[84]: S.No.          0
         Name          0
         Location      0
         Year          0
         Kilometers_Driven 0
         Fuel_Type      0
         Transmission   0
         Owner_Type     0
         Mileage        2
         Engine         46
         Power          175
         Seats          53
         New_price      6247
         Price          1234
         dtype: int64
```

```
In [85]: # Check percentage of missing values
         (data.isnull().sum() / data.shape[0])*100
```

```
Out[85]: S.No.          0.000000
         Name          0.000000
         Location      0.000000
         Year          0.000000
         Kilometers_Driven 0.000000
         Fuel_Type      0.000000
         Transmission   0.000000
         Owner_Type     0.000000
         Mileage        0.027575
         Engine         0.634220
         Power          2.412795
         Seats          0.730732
         New_price      86.129877
         Price          17.013650
         dtype: float64
```

There are missing values in 'Engine', 'Power', 'Seats', 'New\_Price', and 'Price'. \

New\_Price has ~86% missing values, the column may not be useful at all. \ Price has ~17% missing values, this is the dependent variable.

```
In [86]: data.shape
```

```
Out[86]: (7253, 14)
```

The data has 7253 rows and 14 columns.

```
In [87]: # Check for duplicates
         data.duplicated().sum()
```

```
Out[87]: 0
```

There is no duplicated information.

```
In [88]: # Checking descriptive statistics
```

```
data.describe(include = 'all').T
```

Out [88]:

	count	unique	top	freq	mean	std	n
<b>S.No.</b>	7253.0	NaN	NaN	NaN	3626.0	2093.905084	(
<b>Name</b>	7253	2041	Mahindra XUV500 W8 2WD	55	NaN	NaN	Na
<b>Location</b>	7253	11	Mumbai	949	NaN	NaN	Na
<b>Year</b>	7253.0	NaN	NaN	NaN	2013.365366	3.254421	1996
<b>Kilometers_Driven</b>	7253.0	NaN	NaN	NaN	58699.063146	84427.720583	17
<b>Fuel_Type</b>	7253	5	Diesel	3852	NaN	NaN	Na
<b>Transmission</b>	7253	2	Manual	5204	NaN	NaN	Na
<b>Owner_Type</b>	7253	4	First	5952	NaN	NaN	Na
<b>Mileage</b>	7251.0	NaN	NaN	NaN	18.14158	4.562197	(
<b>Engine</b>	7207.0	NaN	NaN	NaN	1616.57347	595.285137	72
<b>Power</b>	7078.0	NaN	NaN	NaN	112.765214	53.493553	34
<b>Seats</b>	7200.0	NaN	NaN	NaN	5.280417	0.809277	2
<b>New_price</b>	1006.0	NaN	NaN	NaN	22.779692	27.759344	3.
<b>Price</b>	6019.0	NaN	NaN	NaN	9.479468	11.187917	0.

### Observations

- There are missing values on some of the columns.
- The dependent variable (Price) has missing values.
- There are two types of variables, numeric and categorical.
- There are extreme values in some columns, like in kilometers driven where one value is of 6,500,000.
- Variables may not follow normal distributions.

## Exploratory Data Analysis

- EDA is an important part of any project involving data.
- It is important to investigate and understand the data better before building a model with it.

In [89]: *# Function to plot a boxplot and a histogram along the same scale*

```
def histogram_boxplot(data, feature, figsize = (12, 7), kde = True, bins = N
    """
```

**Boxplot and histogram combined**

```

data: dataframe
feature: dataframe column
figsize: size of figure (default (12,7))
kde: whether to show density curve (default False)
bins: number of bins for histogram (default None)
"""
f2, (ax_box2, ax_hist2) = plt.subplots(
    nrows = 2,          # Number of rows of the subplot grid = 2
    sharex = True,      # x-axis will be shared among all subplots
    gridspec_kw = {"height_ratios": (0.25, 0.75)},
    figsize = figsize,
)
# Creating the 2 subplots
sns.boxplot(
    data = data, x = feature, ax = ax_box2, showmeans = True, color = "violet"
)
# Boxplot will be created and a star will indicate the mean
sns.histplot(
    data = data, x = feature, kde = kde, ax = ax_hist2, bins = bins, palette = "magma"
)
if bins else sns.histplot(
    data = data, x = feature, kde = kde, ax = ax_hist2
)
# For histogram
ax_hist2.axvline(
    data[feature].mean(), color = "green", linestyle = "--"
)
# Add mean to the histogram
ax_hist2.axvline(
    data[feature].median(), color = "black", linestyle = "-"
)
# Add median to the histogram

```

## Univariate Analysis

```

In [90]: # Create lists with numerical and categorical variable names.
num_variables = ['S.No.', 'Year', 'Kilometers_Driven', 'Mileage', 'Engine', 'Power']
cat_variables = ['Name', 'Location', 'Fuel_Type', 'Transmission', 'Owner_Type']

```

```

In [91]: # Check unique values of Serial Number
data['S.No.'].value_counts()
data.drop('S.No.', axis = 1, inplace = True)


```

All the values in the serial number columns are different, no relevant information for predicting is provided.

```


In [92]: # Histogram and boxplot for year
histogram_boxplot(data, 'Year')

```



The data is a bit skewed to the left, but all values seem reasonable.

```
In [93]: # Histogram and boxplot for Kilometers driven
          histogram_boxplot(data, 'Kilometers_Driven')
```

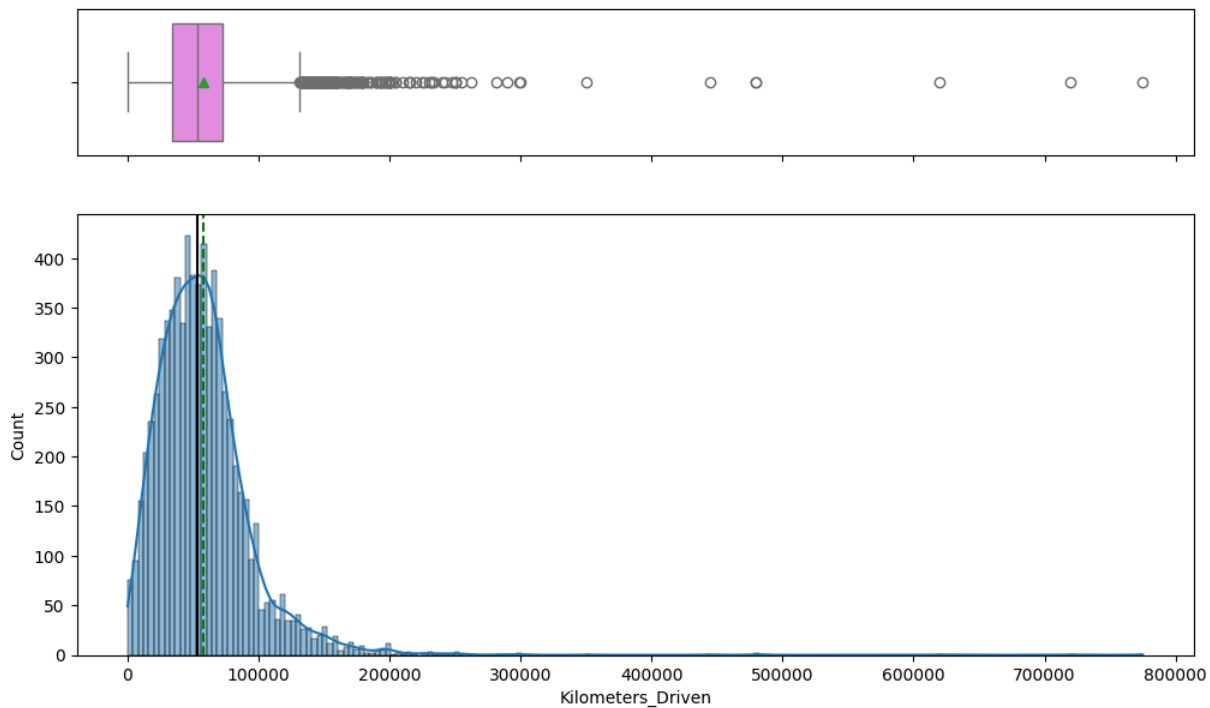


The graph looks highly skewed to the right, it is unreasonable to think that a car still works after 6.5 million kilometers.

```
In [94]: # Drop the row with the extreme value
          data.drop(index = 2328, inplace = True)
```




```
In [95]: # Histogram and boxplot for Kilometers Driven
histogram_boxplot(data, 'Kilometers_Driven' )
```



The data is still skewed to the right, as there are less cars that have a high kilometer count, performing a log or sqrt transform might help normalize data.


```
In [96]: # Transform Kilometers_driven
data['Kilometers_Driven_Transformed'] = np.sqrt(data['Kilometers_Driven'])
```

```
In [97]: # Plot the transformed variable
histogram_boxplot(data, 'Kilometers_Driven_Transformed', kde = True)
```




It is still slightly skewed, but it is better than before.

```
In [98]: # Histogram and boxplot for Mileage
         histogram_boxplot(data, 'Mileage')
```




The 'Mileage' variable has some outliers, however they are not extreme, most of them are proper values, except the 0's, since no car can have 0 mileage.

```
In [99]: # Histogram and boxplot for Engine
         histogram_boxplot(data, 'Engine')
```




There are a few outliers in the data, it is skewed to the right.

```
In [100... # Histogram and boxplot for power
histogram_boxplot(data, 'Power')
```




Another variable that is skewed, we also have some outliers in the upper quartile.

```
In [101... # Histogram and boxplot for seats
histogram_boxplot(data, 'Seats')
```




There is no particular distribution since the seats can only be an integer, all the values are proper values.

```
In [102... # Histogram and boxplot for new price
histogram_boxplot(data, 'New_price')
```



Data is skewed to the right, there are a lot of outliers.


```
In [103... # Histogram and boxplot for price
histogram_boxplot(data, 'Price')
```



This is the dependent variable, it is skewed to the right and has a lot of outliers.  
Performing a log or sqrt transform might be useful.


```
In [104... data['Price_Transformed'] = np.log(data['Price'])
```

```
In [105... # Histogram and boxplot for price
histogram_boxplot(data, 'Price_Transformed')
```



Now the data is more normally distributed

```
In [106... # Plot the countplot for car names
plt.figure(figsize=(14, 8))
sns.countplot(x='Name', data=data, order=data['Name'].value_counts().index)
plt.title('Car names distribution')
plt.ylabel('Count')
plt.xlabel('Car Name')
plt.xticks(rotation = 90)
plt.show()
```




Some models are a lot more popular than others, this can give information about popularity and trends.

```
In [107... top_car_names = data['Name'].value_counts().nlargest(25)

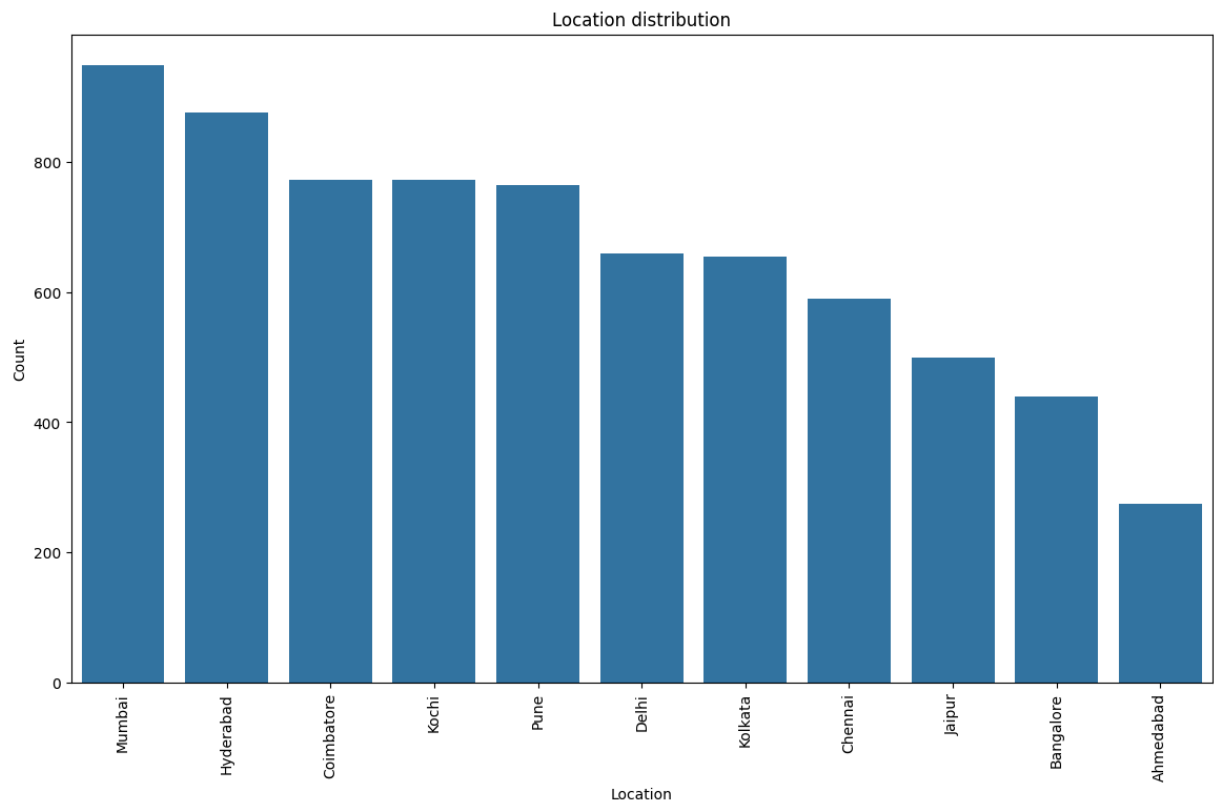
# Plot the countplot for the top 30 car names
plt.figure(figsize=(14, 8))
sns.countplot(x='Name', data=data, order=top_car_names.index)
plt.title('Top 30 Most Frequent Car Names')
plt.ylabel('Count')
```

```
plt.xlabel('Car Name')
plt.xticks(rotation = 90)
plt.show()
```



These are the top 25 models, all of them have a frequency of over 20. We can see that hyundai is quite popular. Splitting this variable into multiple variables like Brand and model might be useful to be able to get more information.

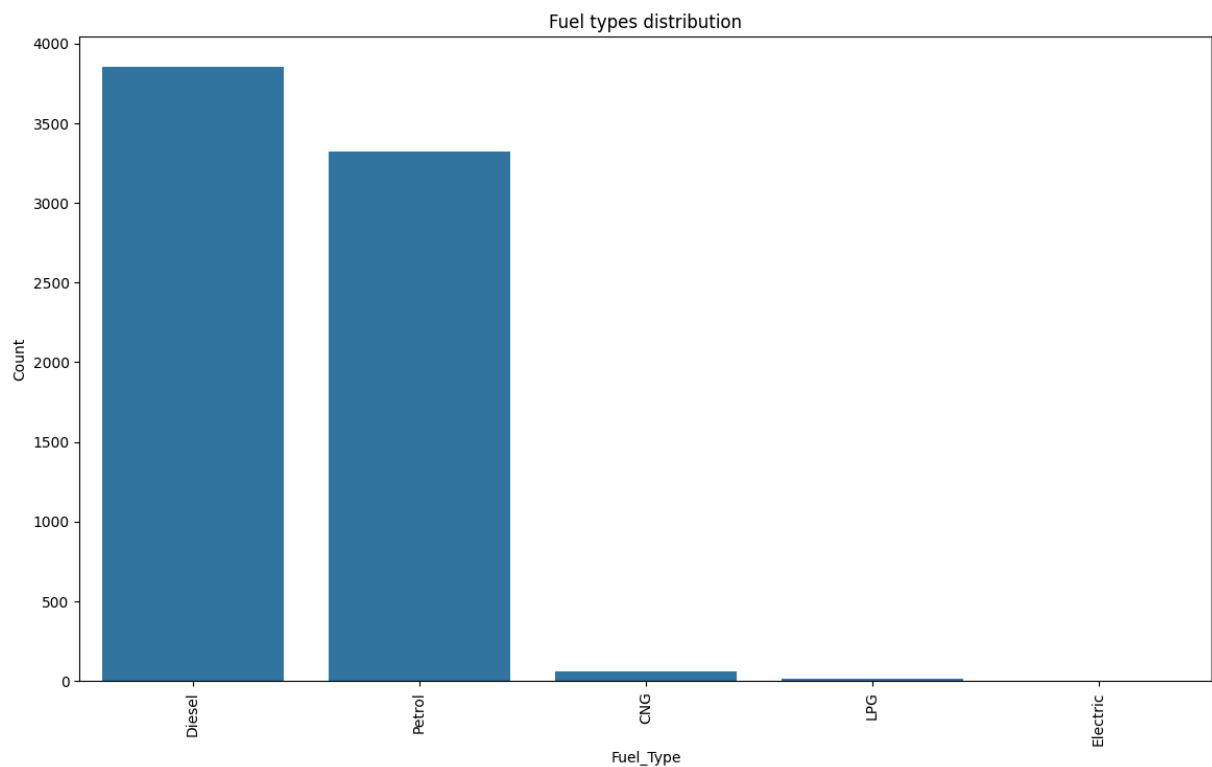
```
In [108... # Plot the countplot for location
plt.figure(figsize=(14, 8))
sns.countplot(x='Location', data=data, order=data['Location'].value_counts())
plt.title('Location distribution')
plt.ylabel('Count')
plt.xlabel('Location')
plt.xticks(rotation = 90)
plt.show()
```



There are not a lot of locations, just 11, with the largest being Mumbai and the lowest being Ahmedabad.


```
In [109... # Plot the countplot for car names
plt.figure(figsize=(14, 8))
sns.countplot(x='Fuel_Type', data=data, order=data['Fuel_Type'].value_counts)
plt.title('Fuel types distribution')
plt.ylabel('Count')
plt.xlabel('Fuel_Type')
plt.xticks(rotation = 90)
plt.show()
```





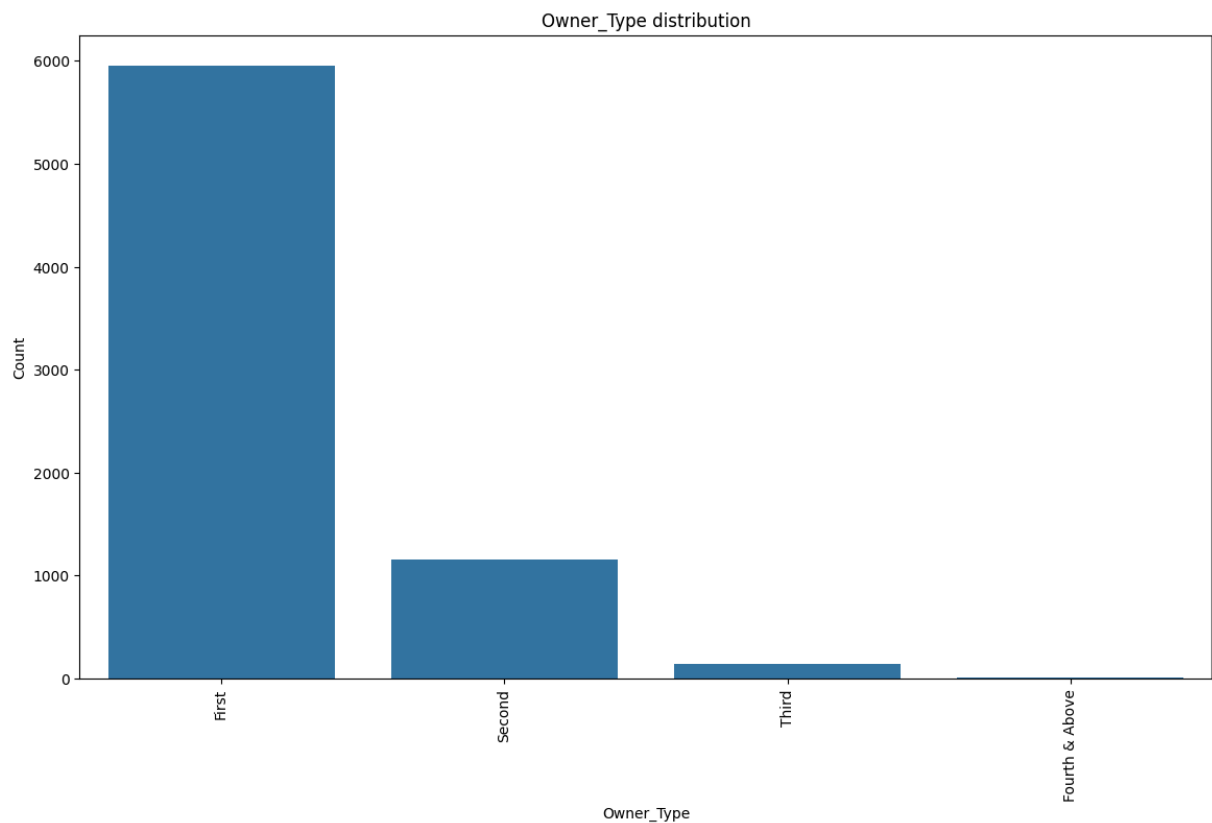
The most common type of fuel is Diesel, the least common is Electric. There are only 5 types of fuel.

```
In [110... # Plot the countplot for car names
plt.figure(figsize=(14, 8))
sns.countplot(x='Transmission', data=data, order=data['Transmission'].value_
plt.title('Transmission distribution')
plt.ylabel('Count')
plt.xlabel('Transmission')
plt.xticks(rotation = 90)
plt.show()
```



There are only two types of transmission, with manual being the most popular.

```
In [111... # Plot the countplot for car names
plt.figure(figsize=(14, 8))
sns.countplot(x='Owner_Type', data=data, order=data['Owner_Type'].value_cour
plt.title('Owner_Type distribution')
plt.ylabel('Count')
plt.xlabel('Owner_Type')
plt.xticks(rotation = 90)
plt.show()
```




There are 4 types of Owner, the most popular being first and the least fourth.

## Bivariate Analysis

```
In [112... num_df = data[['Year', 'Kilometers_Driven_Transformed', 'Mileage', 'Engine', 'Po
plt.figure(figsize = (15,7))
sns.heatmap(num_df.corr(), annot = True, vmin = -1, vmax = 1, fmt = ".2f", c
```

```
Out[112... <Axes: >
```




There are some significant correlations ( $>0.7$  or  $<-0.7$ ) shown in the heatmap.

- Power and Engine (0.86)
- New Price and Engine (0.74)
- Price transformed and Engine (0.69)
- New Price and Power (0.88)
- Price Transformed and Power (0.77)
- New Price and Price Transformed (0.79)

```
In [113...] sns.scatterplot(x = data['Power'], y = data['Engine'])
```


```
Out[113...] <Axes: xlabel='Power', ylabel='Engine'>
```



As the engine's displacement volume increases, so does Power.

```
In [114...] sns.scatterplot(x = data['Engine'], y = data['New_price'])
```


```
Out[114...] <Axes: xlabel='Engine', ylabel='New_price'>
```



As engine's displacement volume increases so does new price, similar thing happens with new price and power.

```
In [115...] sns.scatterplot(x = data['Engine'], y = data['Price_Transformed'])
```


```
Out[115...] <Axes: xlabel='Engine', ylabel='Price_Transformed'>
```



As engine's displacement volume increases, so does price, similar thing happens with price and power.

```
In [116...] sns.scatterplot(x = data['New_price'], y = data['Price_Transformed'])
```

```
Out[116...] <Axes: xlabel='New_price', ylabel='Price_Transformed'>
```




There is correlation between new\_price and price\_transformed. This is expected since expensive new cars can become relatively expensive used cars.

```
In [117...] sns.boxplot(x = data['Price'], y = data['Location'])
```

```
Out[117...] <Axes: xlabel='Price', ylabel='Location'>
```






The distribution of prices may vary a little depending on the location.

```
In [118...] sns.boxplot(x = data['Price'], y = data['Fuel_Type'])
```

```
Out[118...] <Axes: xlabel='Price', ylabel='Fuel_Type'>
```




```
In [119...] data[data['Fuel_Type']=='Electric']
```

```
Out[119...]
      Name  Location  Year  Kilometers_Driven  Fuel_Type  Transmission  Owner_
4446  Mahindra  E Verito  Chennai  2016           50000      Electric      Automatic
      D4
4904  Toyota    Prius    Mumbai  2011           44000      Electric      Automatic
      2009-2016 Z4
```

There are only 2 electric vehicles listed. It seems that vehicles that use diesel are more expensive.

```
In [120...] sns.boxplot(x = data['Price'], y = data['Transmission'])
```


```
Out[120...] <Axes: xlabel='Price', ylabel='Transmission'>
```



Automatic transmission cars are more expensive, as expected.

```
In [121...] sns.boxplot(x = data['Price'], y = data['Owner_Type'])
```

```
Out[121...] <Axes: xlabel='Price', ylabel='Owner_Type'>
```



It makes sense that when the number of previous owners is lower, the price is higher.

## Feature Engineering

```
In [122... # Split the 'Name' column into 'Brand' and 'Model'
data[['Brand', 'Model']] = data['Name'].str.split(n=1, expand=True)

# Display the first few rows of the dataframe to verify the new columns
data.head(7)
```

```
Out [122...      Name      Location  Year  Kilometers_Driven  Fuel_Type  Transmission  Owner_1
0  Maruti      Wagon R LXI CNG      Mumbai  2010           72000         CNG         Manual
1  Hyundai      Creta 1.6 CRDi SX Option      Pune  2015           41000        Diesel         Manual
2  Honda      Jazz V      Chennai  2011           46000        Petrol         Manual
3  Maruti      Ertiga VDI      Chennai  2012           87000        Diesel         Manual
4  Audi A4      New 2.0 TDI Multitronic      Coimbatore  2013           40670        Diesel      Automatic      Sec
5  Hyundai      EON LPG Era Plus Option      Hyderabad  2012           75000         LPG         Manual
6  Nissan      Micra Diesel XV      Jaipur  2013           86999        Diesel         Manual
```

```
In [123... data['Brand'].value_counts()
```

```
Out[123... Brand
Maruti          1444
Hyundai         1340
Honda           743
Toyota          507
Mercedes-Benz   380
Volkswagen      374
Ford            351
Mahindra        331
BMW             311
Audi            285
Tata            228
Skoda           202
Renault         170
Chevrolet       151
Nissan           117
Land            67
Jaguar          48
Fiat            38
Mitsubishi      36
Mini            31
Volvo           28
Porsche         19
Jeep            19
Datsun          17
ISUZU           3
Force           3
Isuzu           2
Bentley         2
Smart           1
Ambassador      1
Lamborghini     1
Hindustan       1
OpelCorsa       1
Name: count, dtype: int64
```

```
In [124... data['Model'].value_counts()
```

```
Out[124... Model
XUV500 W8 2WD          55
Swift VDI              49
Swift Dzire VDI        42
City 1.5 S MT          39
Swift VDI BSIV         37
..
Manza Aura Plus Quadrajet BS IV  1
Indigo eCS LS (TDI) BS-III      1
Grand i10 Era           1
Figo Diesel             1
Elite i20 Magna Plus      1
Name: count, Length: 2041, dtype: int64
```

## Missing value treatment

```
In [125... data.isnull().sum()
```

```
Out[125... Name          0
Location        0
Year            0
Kilometers_Driven  0
Fuel_Type       0
Transmission    0
Owner_Type      0
Mileage         2
Engine          46
Power          175
Seats           53
New_price      6246
Price          1234
Kilometers_Driven_Transformed  0
Price_Transformed  1234
Brand          0
Model          0
dtype: int64
```

### Mileage missing values

```
In [126... # Impute null values with median
data['Mileage'].fillna(data['Mileage'].median(), inplace=True)
```

```
In [127... # Check that the previous method worked
data['Mileage'].isnull().sum()
```

```
Out[127... 0
```

Also treat cases where mileage = 0

```
In [128... # Find rows where 'Mileage' is 0
missing_mileage = data[data['Mileage'] == 0]

# Find rows where 'Mileage' is not 0
not_missing_mileage = data[data['Mileage'] != 0]

# Impute missing 'Mileage' values based on the same model's existing values
for index, row in missing_mileage.iterrows():
    model = row['Model']
    mileage_value = not_missing_mileage[not_missing_mileage['Model'] == model]
    if pd.notnull(mileage_value) and mileage_value != 0:
        data.loc[index, 'Mileage'] = mileage_value

# Fill any remaining missing values with the overall median
median_mileage = data[data['Mileage'] != 0]['Mileage'].median()
data.loc[data['Mileage'] == 0, 'Mileage'] = median_mileage
```

```
In [129... data[data['Mileage'] == 0]
```

Out[129... **Name Location Year Kilometers\_Driven Fuel\_Type Transmission Owner\_Type M**

This ensures that the values were treated properly.

### Engine missing values

```
In [130... # Find rows where 'Engine' is missing
missing_engine = data[data['Engine'].isnull()]

# Find rows where 'Engine' is not missing
not_missing_engine = data[data['Engine'].notnull()]

# Impute missing 'Engine' values based on the same model's existing values
for index, row in missing_engine.iterrows():
    model = row['Model']
    engine_value = not_missing_engine[not_missing_engine['Model'] == model]['Engine'].values[0]
    if pd.notnull(engine_value):
        data.loc[index, 'Engine'] = engine_value

# Fill any remaining missing values with the overall median
median_engine = data['Engine'].median()
data['Engine'].fillna(median_engine, inplace=True)
```

```
In [131... data['Engine'].isnull().sum()
```

Out[131... 0

Now there are no more missing Engine values.

### Power missing values

```
In [132... # Find rows where 'Power' is missing
missing_power = data[data['Power'].isnull()]

# Find rows where 'Power' is not missing
not_missing_power = data[data['Power'].notnull()]

# Impute missing 'Power' values based on the same model's existing values
for index, row in missing_power.iterrows():
    model = row['Model']
    power_value = not_missing_power[not_missing_power['Model'] == model]['Power'].values[0]
    if pd.notnull(power_value):
        data.loc[index, 'Power'] = power_value

# Fill any remaining missing values with the overall median
median_power = data['Power'].median()
data['Power'].fillna(median_power, inplace=True)
```

```
In [133... data['Power'].isnull().sum()
```

Out[133... 0

No more missing values in 'Power'

## Seats missing values

```
In [134... # Find rows where 'Seats' is missing
missing_seats = data[data['Seats'].isnull()]

# Find rows where 'Seats' is not missing
not_missing_seats = data[data['Seats'].notnull()]

# Impute missing 'Seats' values based on the same model's existing values
for index, row in missing_seats.iterrows():
    model = row['Model']
    seats_value = not_missing_seats[not_missing_seats['Model'] == model]['Se
    if pd.notnull(seats_value):
        data.loc[index, 'Seats'] = seats_value

# Fill any remaining missing values with the overall median
median_seats = data['Seats'].median()
data['Seats'].fillna(median_seats, inplace=True)
```

```
In [135... data.isnull().sum()
```

```
Out[135... Name                                0
Location                                0
Year                                    0
Kilometers_Driven                       0
Fuel_Type                               0
Transmission                            0
Owner_Type                              0
Mileage                                  0
Engine                                  0
Power                                    0
Seats                                    0
New_price                               6246
Price                                    1234
Kilometers_Driven_Transformed            0
Price_Transformed                        1234
Brand                                    0
Model                                    0
dtype: int64
```

## New\_price missing values

Since there are a lot of new\_price missing values, comparing models will not work on the majority of cases, so in order to get a more precise replacement of null values, a comparison between a few attributes might help.

```
In [136... # Find rows where 'New_price' is missing
missing_new_price = data[data['New_price'].isnull()]

# Find rows where 'New_price' is not missing
```



```

not_missing_new_price = data[data['New_price'].notnull()]

# Step 1: Model-specific median
for index, row in missing_new_price.iterrows():
    model = row['Model']
    new_price_value = not_missing_new_price[not_missing_new_price['Model'] =
    if pd.notnull(new_price_value):
        data.loc[index, 'New_price'] = new_price_value

# Update missing_new_price after first step
missing_new_price = data[data['New_price'].isnull()]

# Step 2: Brand and Year
for index, row in missing_new_price.iterrows():
    brand = row['Brand']
    year = row['Year']
    new_price_value = not_missing_new_price[(not_missing_new_price['Brand']
    if pd.notnull(new_price_value):
        data.loc[index, 'New_price'] = new_price_value

# Update missing_new_price after second step
missing_new_price = data[data['New_price'].isnull()]

# Step 3: Brand and Fuel_Type
for index, row in missing_new_price.iterrows():
    brand = row['Brand']
    fuel_type = row['Fuel_Type']
    new_price_value = not_missing_new_price[(not_missing_new_price['Brand']
    if pd.notnull(new_price_value):
        data.loc[index, 'New_price'] = new_price_value

# Update missing_new_price after third step
missing_new_price = data[data['New_price'].isnull()]

# Step 4: Overall Brand
for index, row in missing_new_price.iterrows():
    brand = row['Brand']
    new_price_value = not_missing_new_price[not_missing_new_price['Brand'] =
    if pd.notnull(new_price_value):
        data.loc[index, 'New_price'] = new_price_value

# Verify the imputation
missing_new_price = data[data['New_price'].isnull()]

```

In [137... data['New\_price'].isnull().sum()

Out[137... 162

There are still 162 missing values, since these are very few compared to the initial 6+ thousand, the overall median will be used.

In [138... data['New\_price'].fillna(data['New\_price'].median(), inplace=True)

In [139... data['New\_price'].isnull().sum()

Out[139... 0

There are no more missing values.

## Price missing values

Since this is the dependent variable, imputing may affect the quality of the models.

```
In [140... # Drop rows where Price is missing
data.dropna(subset=['Price'], inplace = True)
```

```
In [141... data.isnull().sum()
```

```
Out[141... Name                                0
Location                                0
Year                                    0
Kilometers_Driven                      0
Fuel_Type                             0
Transmission                          0
Owner_Type                            0
Mileage                               0
Engine                                0
Power                                 0
Seats                                 0
New_price                             0
Price                                 0
Kilometers_Driven_Transformed          0
Price_Transformed                     0
Brand                                  0
Model                                  0
dtype: int64
```

There are no more missing values within the dataset.

## Important Insights from EDA and Data Preprocessing

- There was a square root transformation made on Kilometers\_Driven and a Log transformation made to the dependent variable.
- Missing values were imputed based on other features like Model, Brand and/or Fuel type.
- Feature 'Name' was separated into two, Brand and Model.

## Building Various Models

1. What we want to predict is the "Price". We will use the normalized version 'price\_log' for modeling.

2. Before we proceed to the model, we'll have to encode categorical features. We will drop categorical features like Name.
3. We'll split the data into train and test, to be able to evaluate the model that we build on the train data.
4. Build Regression models using train data.
5. Evaluate the model performance.

## Split the Data

- Step1: Separating the independent variables (X) and the dependent variable (y).
- Step2: Encode the categorical variables in X using pd.dummies.
- Step3: Split the data into train and test using train\_test\_split.

## Linear regression

```
In [142... # Function to check VIF
check_vif = data.select_dtypes(include=['number']).drop(columns = ['Price_Tr
def checking_vif(train):
    vif = pd.DataFrame()
    vif["feature"] = train.columns

    # Calculating VIF for each feature
    vif["VIF"] = [
        variance_inflation_factor(train.values, i) for i in range(len(train.
    ]
    return vif

print(checking_vif(check_vif))
```

	feature	VIF
0	Mileage	7.024115
1	Power	8.234872
2	New_price	3.130134
3	Kilometers_Driven_Transformed	8.439496

After checking the collinearity between variables, these are the only ones remaining to get VIF scores under 10.

```
In [143... # Take dependent and independent variables
X = data.drop(['Name', 'Price', 'Price_Transformed', 'Kilometers_Driven', 'Y
X = pd.get_dummies(X, drop_first=True)
X = sm.add_constant(X) # Add constant

Y = data['Price_Transformed']

# Split the data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, ran

# Convert all boolean columns to integers
bool_columns = X_train.select_dtypes(include='bool').columns
```

```
for col in bool_columns:
    X_train[col] = X_train[col].astype(int)
```

```
In [144... # Model Performance on test and train data
def model_performance(olsmodel, x_train, x_test, y_train, y_test):

    # In-sample Prediction
    y_pred_train = olsmodel.predict(x_train)
    y_observed_train = y_train

    # Prediction on test data
    y_pred_test = olsmodel.predict(x_test)
    y_observed_test = y_test

    print(
        pd.DataFrame(
            {
                "Data": ["Train", "Test"],
                "RMSE": [
                    np.sqrt(mean_squared_error(y_observed_train, y_pred_train)),
                    np.sqrt(mean_squared_error(y_observed_test, y_pred_test))
                ],
                "MAE": [
                    mean_absolute_error(y_observed_train, y_pred_train),
                    mean_absolute_error(y_observed_test, y_pred_test)
                ],
                "MAPE": [
                    mean_absolute_percentage_error(y_observed_train, y_pred_train),
                    mean_absolute_percentage_error(y_observed_test, y_pred_test)
                ],
                "r2": [
                    r2_score(y_observed_train, y_pred_train),
                    r2_score(y_observed_test, y_pred_test)
                ]
            }
        )
    )
```

```
In [145... # Train model
linear1 = sm.OLS(Y_train, X_train).fit()
model_performance(linear1, X_train, X_test, Y_train, Y_test)
```

	Data	RMSE	MAE	MAPE	r2
0	Train	0.154192	0.100564	2.535729e+12	0.968666
1	Test	0.384684	0.238129	2.106962e+12	0.808400

Removing the variables that have a p-value higher than 0.05 may improve the model's performance.

```
In [146... # Extract p-values
p_values = linear1.pvalues

# Identify columns with p-values > 0.05 (excluding the constant)
```

```

columns_to_drop = p_values[p_values > 0.05].index
columns_to_drop = columns_to_drop[columns_to_drop != 'const']

# Drop these columns from the dataframe
X = X.drop(columns=columns_to_drop)

```

```

In [147... X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.3, ran
# Convert all boolean columns to integers
bool_columns = X_train.select_dtypes(include='bool').columns

for col in bool_columns:
    X_train[col] = X_train[col].astype(int)

# Refit the model with remaining columns
linear2 = sm.OLS(Y_train, X_train).fit()
model_performance(linear2, X_train, X_test, Y_train, Y_test)
#print(linear2.summary())

```

	Data	RMSE	MAE	MAPE	r2
0	Train	0.180648	0.127895	2.713904e+12	0.956992
1	Test	0.321485	0.222521	2.330986e+12	0.866184

```

In [148... # Repeat the steps
# Extract p-values
p_values = linear2.pvalues

# Identify columns with p-values > 0.05 (excluding the constant)
columns_to_drop = p_values[p_values > 0.05].index
columns_to_drop = columns_to_drop[columns_to_drop != 'const']

# Drop these columns from the dataframe
X = X.drop(columns=columns_to_drop)

```

```

In [149... X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33 , ran
# Convert all boolean columns to integers
bool_columns = X_train.select_dtypes(include='bool').columns

for col in bool_columns:
    X_train[col] = X_train[col].astype(int)

# Refit the model with remaining columns
linear3 = sm.OLS(Y_train, X_train).fit()
model_performance(linear3, X_train, X_test, Y_train, Y_test)
#print(linear3.summary())

```

	Data	RMSE	MAE	MAPE	r2
0	Train	0.228469	0.156834	1.542889e+12	0.931519
1	Test	0.313798	0.215494	6.489165e+12	0.871222

## Check assumptions of linear regression

### Mean of residuals = 0

```

In [150... residuals = linear3.resid

```

```
np.mean(residuals)
```

```
Out[150...] 1.3645059509100869e-14
```

Mean of residuals is almost equal to 0.

## Homoscedasticity

- Residuals must be symmetrically distributed across the regression line.
- Goldfeldquandt test with  $\alpha = 0.05$
- Null hypotheses: Residuals are homoscedastic.
- Alternate hypotheses: Residuals are heteroscedastic.

```
In [151...] # Perform test and display results
name = ["F statistic", "p-value"]

test = sms.het_goldfeldquandt(Y_train, X_train)

lzip(name, test)
```

```
Out[151...] [('F statistic', 1.037947472721264), ('p-value', 0.23744596342610716)]
```


Since the p-value is greater than 0.05, the assumption holds.

## Linearity of variables

```
In [152...] # Predicted values
fitted = linear3.fittedvalues

sns.residplot(x = fitted, y = residuals)
plt.xlabel("Fitted Values")
plt.ylabel("Residual")
plt.title("Residual PLOT")

plt.show()
```




The residuals are randomly and uniformly scattered along the x axis, they do not form any pattern or follow any trend.

### Normality of error terms

```
In [153... # Plot histogram to see distribution of residuals  
sns.histplot(residuals, kde = True)
```


```
Out[153... <Axes: ylabel='Count'>
```



```
In [154... # Q-Q plot to confirm normality
stats.probplot(residuals, dist = "norm", plot = pylab)

plt.show()
```





The residuals follow a fairly normal distribution.

```
In [155... # Function to perform cross validation
def cross_validate_sm_ols(X, y, k=10):
    "Perform cross-validation, takes the values of the dependent variables a

    kf = KFold(n_splits=k, shuffle=True, random_state=1)
    r2_scores = []
    rmse_scores = []

    for train_index, test_index in kf.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        model = sm.OLS(y_train, X_train).fit()
        y_pred = model.predict(X_test)

        r2_scores.append(r2_score(y_test, y_pred))
        rmse_scores.append(mean_squared_error(y_test, y_pred, squared=False))

    mean_r2 = np.mean(r2_scores)
    std_r2 = np.std(r2_scores)
    mean_rmse = np.mean(rmse_scores)
    std_rmse = np.std(rmse_scores)

    return mean_r2, std_r2, mean_rmse, std_rmse
```

```
In [156... # Use cross validation function for different values of k
results = []
for k in range(2, 10):
    for col in bool_columns:
        X[col] = X[col].astype(int)
    mean_r2, std_r2, mean_mrse, std_mrse = cross_validate_sm_ols(X.values, Y
    results.append((k, mean_r2, 2 * std_r2, mean_mrse, 2 * std_mrse))

results_df = pd.DataFrame(results, columns=['k', 'R-squared', ' +/-', 'MRSE']
print(results_df)
```

	k	R-squared	+/-	MRSE	+/-
0	2	0.856162	0.000409	0.331282	0.004762
1	3	0.866473	0.005125	0.318974	0.008838
2	4	0.870237	0.017880	0.314437	0.028861
3	5	0.873612	0.016838	0.310348	0.036898
4	6	0.874409	0.019001	0.309261	0.032835
5	7	0.876201	0.017040	0.307160	0.031769
6	8	0.875481	0.023062	0.307739	0.034164
7	9	0.876509	0.033370	0.306183	0.049714

Splitting the data in different ratios slightly affect the performance of the model. Overall, the model is able to explain ~87% of the variation.

## Ridge regression

- Now ridge regression will be used to create a model to predict the values. -Ridge regression is used instead of Lasso regression due to the correlation between some variables, Lasso regression is able to make coefficients zero while Ridgre regression is not able to do that.

```
In [168... # Take dependent and independent variables
X = data.drop(['Name', 'Price', 'Price_Transformed', 'Kilometers_Driven'], a
X = pd.get_dummies(X, drop_first=True)

Y = data['Price_Transformed']

# Split the data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, ran

# Convert all boolean columns to integers
bool_columns = X_train.select_dtypes(include='bool').columns

for col in bool_columns:
    X_train[col] = X_train[col].astype(int)
```

```
In [169... # Check the best alpha
params = {'alpha': [0.001, 0.1, 0.2, 0.5, 0.9, 1, 2, 5, 8, 10]}
folds = KFold(n_splits = 10, shuffle = True, random_state = 1)

ridge = Ridge()
```

```
ridge_cv = GridSearchCV(estimator=ridge, param_grid=params, cv=folds, scoring='neg_mean_squared_error')
ridge_cv.fit(X_train, Y_train)
```

Out[169]...

```
GridSearchCV
  estimator: Ridge
    Ridge
```

In [171]...

```
# View the result
ridge_cv.best_params_
```

Out[171]...

```
{'alpha': 0.5}
```

In [172]...

```
# Build the model with the right alpha
ridge_model = Ridge(alpha = 0.5)
ridge_model.fit(X_train, Y_train)
```

Out[172]...

```
Ridge
Ridge(alpha=0.5)
```

In [173]...

```
# Evaluate model performance
model_performance(ridge_model, X_train, X_test, Y_train, Y_test)
```

	Data	RMSE	MAE	MAPE	r2
0	Train	0.123351	0.088638	2.211775e+12	0.979938
1	Test	0.188580	0.128700	1.559865e+12	0.954479

In [174]...

```
# Function to cross validate Ridge
def cross_validate_ride(X, y, alpha=0.5, k=10):
    """
    Perform cross-validation for Ridge regression.
    X: Features
    y: Target variable
    alpha: Regularization strength
    k: Number of folds for cross-validation
    """
    kf = KFold(n_splits=k, shuffle=True, random_state=1)
    r2_scores = []
    rmse_scores = []

    for train_index, test_index in kf.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        model = Ridge(alpha=alpha)
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)

        r2_scores.append(r2_score(y_test, y_pred))
        rmse_scores.append(mean_squared_error(y_test, y_pred, squared=False))

    mean_r2 = np.mean(r2_scores)
```

```
std_r2 = np.std(r2_scores)
mean_rmse = np.mean(rmse_scores)
std_rmse = np.std(rmse_scores)

return mean_r2, std_r2, mean_rmse, std_rmse
```

```
In [175... # Use cross validation function for different values of k
results = []
for k in range(2, 10):
    for col in bool_columns:
        X[col] = X[col].astype(int)
        mean_r2, std_r2, mean_rmse, std_rmse = cross_validate_ridge(X.values, Y,
        results.append((k, mean_r2, 2 * std_r2, mean_rmse, 2 * std_rmse))

results_df = pd.DataFrame(results, columns=['k', 'R-squared', ' +/-', 'MRSE']
print(results_df)
```

	k	R-squared	+/-	MRSE	+/-
0	2	0.944542	0.008589	0.205614	0.019188
1	3	0.946116	0.013123	0.202461	0.030702
2	4	0.947002	0.016497	0.200739	0.038563
3	5	0.948281	0.014026	0.198236	0.034635
4	6	0.948834	0.015132	0.197189	0.036058
5	7	0.948941	0.025387	0.196107	0.048159
6	8	0.949121	0.023119	0.196058	0.046776
7	9	0.949356	0.028700	0.195112	0.056382

- This model has a better performance than the Linear regression model made previously.
- It is able to consistently explain about ~94% - 95% of variance found in the target variable
- Thus far, this is the best model

## Observations

- The optimized value of alpha was of 0.5
- This model performs better than the simple linear regression since Ridge regression is more robust when variables have collinearity.
- The ridge regression model was able to explain about ~94% of the variation.

## Hyperparameter Tuning: Decision Tree

```
In [176... # Check the best parameters
params = {
    'max_depth': [None, 10, 20, 30, 40, 50],
    'min_samples_split': [2, 5, 10, 20],
    'min_samples_leaf': [1, 2, 5, 10]
}
folds = KFold(n_splits = 10, shuffle = True, random_state = 1)

decision_tree = DecisionTreeRegressor()
```

```
decision_tree_cv = GridSearchCV(estimator=decision_tree, param_grid=params,
decision_tree_cv.fit(X_train, Y_train)
```

Out[176... **GridSearchCV**

- **estimator: DecisionTreeRegressor**
  - DecisionTreeRegressor

In [177... *# show best parameters*  
decision\_tree\_cv.best\_params\_

Out[177... {'max\_depth': 50, 'min\_samples\_leaf': 1, 'min\_samples\_split': 20}

Now, use those parameters to build a model

In [184... *# Build model with best parameters*  
tree\_model = DecisionTreeRegressor(max\_depth=50, min\_samples\_leaf=1, min\_sam  
tree\_model.fit(X\_train, Y\_train)

Out[184... **DecisionTreeRegressor**  
DecisionTreeRegressor(max\_depth=50, min\_samples\_split=20)

In [185... *# Evaluate model performance*  
model\_performance(tree\_model, X\_train, X\_test, Y\_train, Y\_test)

	Data	RMSE	MAE	MAPE	r2
0	Train	0.147550	0.104590	2.960163e+12	0.971294
1	Test	0.236967	0.172893	1.657659e+12	0.928122

This model is able to explain around 93% of variance in the target variable.

In [186... *# Function to cross validate Decision Tree*

```
def cross_validate_tree(X, y, k=10, max_depth = 20, min_samples_leaf = 1, mi
    """
    Perform cross-validation for Ridge regression.
    X: Features
    y: Target variable
    alpha: Regularization strength
    k: Number of folds for cross-validation
    """
    kf = KFold(n_splits=k, shuffle=True, random_state=1)
    r2_scores = []
    rmse_scores = []

    for train_index, test_index in kf.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        model = DecisionTreeRegressor(max_depth = max_depth, min_samples_lea
        model.fit(X_train, y_train)
```

```

y_pred = model.predict(X_test)

r2_scores.append(r2_score(y_test, y_pred))
rmse_scores.append(mean_squared_error(y_test, y_pred, squared=False))

mean_r2 = np.mean(r2_scores)
std_r2 = np.std(r2_scores)
mean_rmse = np.mean(rmse_scores)
std_rmse = np.std(rmse_scores)

return mean_r2, std_r2, mean_rmse, std_rmse

```

```

In [187... # Use cross validation function for different values of k
results = []
for k in range(2, 10):
    mean_r2, std_r2, mean_rmse, std_rmse = cross_validate_tree(X.values, Y.v
    results.append((k, mean_r2, 2 * std_r2, mean_rmse, 2 * std_rmse))

results_df = pd.DataFrame(results, columns=['k', 'R-squared', ' +/-', 'MRSE']
print(results_df)

```

	k	R-squared	+/-	MRSE	+/-
0	2	0.898688	0.008460	0.278016	0.016002
1	3	0.908661	0.011952	0.263781	0.023395
2	4	0.910075	0.019965	0.261684	0.038314
3	5	0.910831	0.020143	0.260205	0.033850
4	6	0.910998	0.015782	0.260228	0.029504
5	7	0.912171	0.031418	0.257922	0.050143
6	8	0.909073	0.032058	0.262560	0.054154
7	9	0.913745	0.032252	0.255585	0.053398

This model is better than the simple Linear regression model, but it is not better than the Ridge Regression model.

## Feature Importance

```

In [188... # Access feature importances
feature_importances = tree_model.feature_importances_

# Create a DataFrame for better readability
features = X.columns
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': feature_importances
})

# Sort the DataFrame by importance
importance_df = importance_df.sort_values(by='Importance', ascending=False)


# Select the top 20 features
top_20_features = importance_df.head(20)

# Display the DataFrame
print(top_20_features)

```

```
# Plot the feature importances of the top 20 features
plt.figure(figsize=(12, 8))
plt.barh(top_20_features['Feature'], top_20_features['Importance'])
plt.xlabel('Feature Importance')
plt.ylabel('Feature')
plt.title('Top 20 Feature Importances in Decision Tree')
plt.gca().invert_yaxis()
plt.show()
```

	Feature	Importance
3	Power	0.604303
0	Year	0.231673
5	New_price	0.086054
2	Engine	0.025477
6	Kilometers_Driven_Transformed	0.005342
1	Mileage	0.005087
14	Location_Kolkata	0.004191
51	Brand_Tata	0.003903
52	Brand_Toyota	0.003087
352	Model_Cayenne Base	0.002901
11	Location_Hyderabad	0.002753
9	Location_Coimbatore	0.001609
32	Brand_Ford	0.001552
4	Seats	0.001341
7	Location_Bangalore	0.001261
28	Brand_Chevrolet	0.000846
1120	Model_New Safari DICOR 2.2 VX 4x2	0.000712
17	Fuel_Type_Diesel	0.000685
10	Location_Delhi	0.000647
27	Brand_Bentley	0.000633



In this plot, it is shown that the most important features are:

- Power
- Year

- New\_price
- Engine

## Hyperparameter Tuning: Random Forest

```
In [189... # Check the best parameters
params = {
    'n_estimators': [200, 300],
    'max_depth': [10, 20],
    'min_samples_split': [ 5,10],
    'min_samples_leaf': [ 2, 4]
}

folds = KFold(n_splits = 10, shuffle = True, random_state = 1)

forest = RandomForestRegressor()

forest_cv = GridSearchCV(estimator=forest, param_grid=params, cv=5, scoring=
forest_cv.fit(X_train, Y_train)
```

```
Out[189...
└─ GridSearchCV
  └─ estimator: RandomForestRegressor
    └─ RandomForestRegressor
```

```
In [190... # Display best parameters
forest_cv.best_params_
```

```
Out[190... {'max_depth': 20,
            'min_samples_leaf': 2,
            'min_samples_split': 5,
            'n_estimators': 200}
```

```
In [192... # Build model using the previous parameters
forest_model_1 = RandomForestRegressor(max_depth= 20,
min_samples_leaf= 2,
min_samples_split= 5,
n_estimators= 200)
forest_model_1.fit(X_train,Y_train)
```

```
Out[192...
└─ RandomForestRegressor
  RandomForestRegressor(max_depth=20, min_samples_leaf=2, min_samples_
_split=5,
                        n_estimators=200)
```

```
In [193... # Evaluate model
model_performance(forest_model_1, X_train, X_test, Y_train, Y_test)
```



	Data	RMSE	MAE	MAPE	r2
0	Train	0.117591	0.076429	1.388294e+12	0.981768
1	Test	0.194037	0.138088	1.212597e+12	0.951807

```
In [194... # Compare with default parameters
forest_model_2 = RandomForestRegressor()
forest_model_2.fit(X_train, Y_train)
```

```
Out[194... ▼ RandomForestRegressor
RandomForestRegressor()
```

```
In [195... # Evaluate model
model_performance(forest_model_2, X_train, X_test, Y_train, Y_test)
```

	Data	RMSE	MAE	MAPE	r2
0	Train	0.079750	0.053362	1.095007e+12	0.991614
1	Test	0.189722	0.133691	1.523558e+12	0.953926

```
In [197... # Get parameters
forest_model_2.get_params()
```

```
Out[197... {'bootstrap': True,
'ccp_alpha': 0.0,
'criterion': 'squared_error',
'max_depth': None,
'max_features': 1.0,
'max_leaf_nodes': None,
'max_samples': None,
'min_impurity_decrease': 0.0,
'min_samples_leaf': 1,
'min_samples_split': 2,
'min_weight_fraction_leaf': 0.0,
'n_estimators': 100,
'n_jobs': None,
'oob_score': False,
'random_state': None,
'verbose': 0,
'warm_start': False}
```

## Observations

- Default parameters have a better score

```
In [198... # Function to cross validate Decision Tree
def cross_validate_forest(X, y, k=10, max_depth = None, min_samples_leaf = 1
    """
    Perform cross-validation for Ridge regression.
    X: Features
    y: Target variable
    alpha: Regularization strength
    k: Number of folds for cross-validation
    """
    kf = KFold(n_splits=k, shuffle=True, random_state=1)
```

```

r2_scores = []
rmse_scores = []

for train_index, test_index in kf.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    model = RandomForestRegressor(max_depth = max_depth, min_samples_leaf = min_samples_leaf)
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    r2_scores.append(r2_score(y_test, y_pred))
    rmse_scores.append(mean_squared_error(y_test, y_pred, squared=False))

mean_r2 = np.mean(r2_scores)
std_r2 = np.std(r2_scores)
mean_rmse = np.mean(rmse_scores)
std_rmse = np.std(rmse_scores)

return mean_r2, std_r2, mean_rmse, std_rmse

```

```

In [199]: # Use cross validation function for different values of k
results = []
for k in range(2, 10):
    mean_r2, std_r2, mean_rmse, std_rmse = cross_validate_forest(X.values, Y.values, k=k)
    results.append((k, mean_r2, 2 * std_r2, mean_rmse, 2 * std_rmse))

results_df = pd.DataFrame(results, columns=['k', 'R-squared', '+/-', 'MRSE'])
print(results_df)

```

	k	R-squared	+/-	MRSE	+/-
0	2	0.934036	0.005822	0.224330	0.013445
1	3	0.939588	0.011709	0.214475	0.027168
2	4	0.941751	0.013736	0.210611	0.032783
3	5	0.943286	0.015306	0.207596	0.036372
4	6	0.943299	0.016442	0.207527	0.036172
5	7	0.942712	0.025025	0.208009	0.046738
6	8	0.943359	0.019398	0.207269	0.040944
7	9	0.943616	0.031622	0.205814	0.059160

This model is able to account for ~94% of variation.

## Feature Importance

```

In [200]: # Access feature importances
feature_importances = forest_model_2.feature_importances_

# Create a DataFrame for better readability
features = X.columns
importance_df = pd.DataFrame({
    'Feature': features,
    'Importance': feature_importances
})

# Sort the DataFrame by importance

```

```

importance_df = importance_df.sort_values(by='Importance', ascending=False)

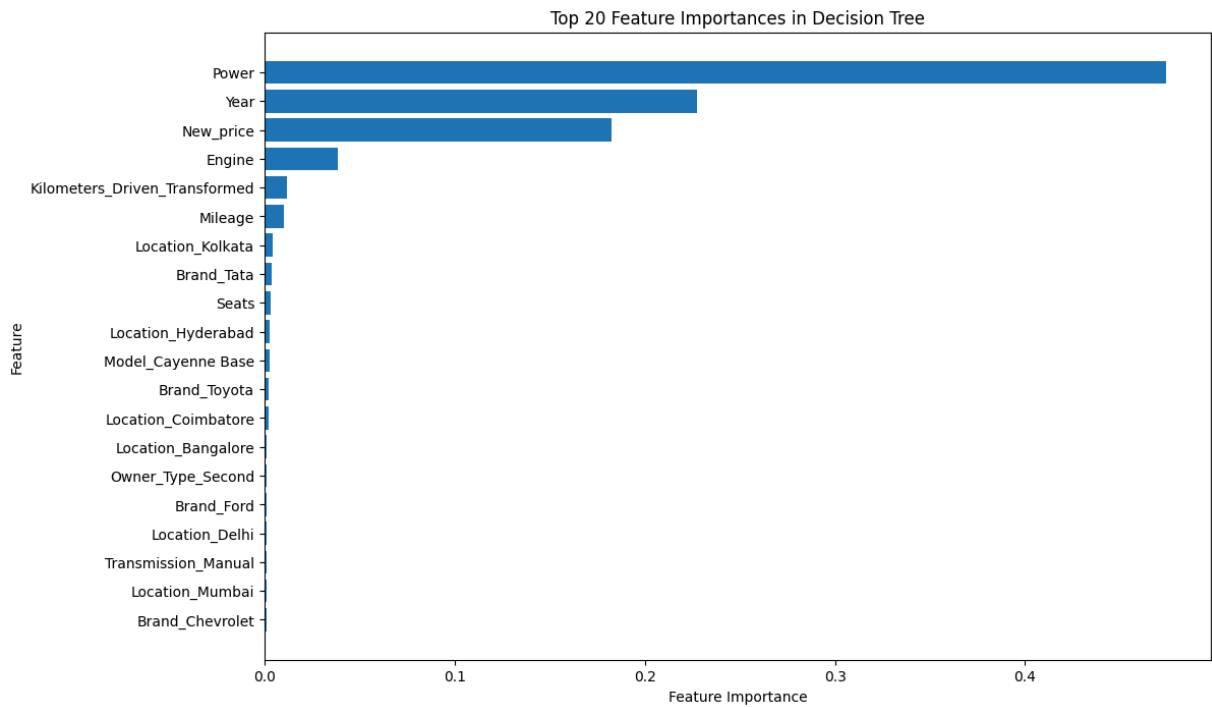
# Select the top 20 features
top_20_features = importance_df.head(20)

# Display the DataFrame
print(top_20_features)

# Plot the feature importances of the top 20 features
plt.figure(figsize=(12, 8))
plt.barh(top_20_features['Feature'], top_20_features['Importance'])
plt.xlabel('Feature Importance')
plt.ylabel('Feature')
plt.title('Top 20 Feature Importances in Decision Tree')
plt.gca().invert_yaxis()
plt.show()

```

	Feature	Importance
3	Power	0.473914
0	Year	0.227295
5	New_price	0.182488
2	Engine	0.038705
6	Kilometers_Driven_Transformed	0.011692
1	Mileage	0.010077
14	Location_Kolkata	0.004439
51	Brand_Tata	0.003874
4	Seats	0.003037
11	Location_Hyderabad	0.002861
352	Model_Cayenne Base	0.002683
52	Brand_Toyota	0.002424
9	Location_Coimbatore	0.001981
7	Location_Bangalore	0.001275
23	Owner_Type_Second	0.001268
32	Brand_Ford	0.001238
10	Location_Delhi	0.001128
21	Transmission_Manual	0.001125
15	Location_Mumbai	0.001082
28	Brand_Chevrolet	0.000965



The most important features for the random forest regression are:

- Power
- Year
- New price
- Engine

These are the same as in the decision tree regression.

## Conclusions and Recommendations

### 1. Comparison of various techniques and their relative performance based on chosen Metric (Measure of success):

The techniques have similar performances. The primary metric used for evaluating performance was  $r^2$ . The ranking is the following

- 1. Ridge Regression  $r^2$  of ~95%
- 2. Random forest  $r^2$  of ~94%
- 3. Decision Tree  $r^2$  of ~91%
- 4. Linear regression  $r^2$  of ~87%

### 2. Refined insights:

All the models are able to make predictions of the price of cars. However, the best models are Ridge Regression and Random forest regression.

The reason is that there is multicollinearity between variables and these two models are robust to these conditions.

The data had to be processed so it would be more effective for training models, a couple transformations were made (sqrt and log) to achieve normality in skewed variables.

### **3. Proposal for the final solution design:**

The best model to adopt is the random forest regression, even though it is slightly worse than the ridge regression. It is far more customizable in the sense that there are more hyper parameters to adjust, the downside is that it takes more resources and time.

This model is also more interpretable since it has the 'Importance' feature, in the analysis made the top 4 features were Power, Engine, Year and New price.

#### **Benefits**

- Improved pricing accuracy.
- Increased sales.
- Improved customer satisfaction.

#### **Costs**

- Initial setup costs for model deployment and infrastructure.
- Ongoing costs for data maintenance and model retraining.